

Demo Speaker Notes — First 3 Sections

Lowe's AI Sales Targeting · Streamlit Walkthrough

Capstone 8 · Group 8 · IIM Calcutta APAL02 | First ~2 minutes of the demo

Section	Time
0 • Opening + The 94.7% hero banner	0:00 – 0:20
1 • Pick a real store — Wolfridge, Texas (a 'Best Save')	0:20 – 1:15
2 • Click EXPLAIN — the GenAI explainer	1:15 – 1:50
Bonus • How the GenAI layer actually works (for Q&A)	+1 page

0 • Opening + The 94.7% Hero Banner

0:00 – 0:20 (~20 seconds)

WHAT TO DO

- Open the browser tab that already has the dashboard loaded (do this BEFORE the demo starts).
- Scroll to the very top so the big '94.7%' number is the first thing they see.

WHAT TO SAY

"This is the dashboard our planner would actually use every week. The first thing you see is the headline result — 94.7% of stores now get an accurate target, within plus-or-minus five percent. So instead of reviewing all 1,727 stores manually like they do today, the team only needs to look at 91. That is a 95% cut in their workload."

TIP Pre-open the dashboard before the demo. Don't switch tabs live. Have it ready.

NEXT Scroll down to the section titled 'Try It: Pick a Store'.

1 • Pick a Real Store — Wolfridge, Texas

0:20 – 1:15 (~55 seconds)

WHAT TO DO

- In the dropdown, start typing 'Wolfridge' or 'S01155'. The store will appear in the list.
- Click it. Wait for the three cards (Actual, Peanut-butter, Model) to show on screen.

WHAT TO SAY

"Let me show you what this means for a real store. I picked Wolfridge, Texas. Look at the three numbers on screen. The actual sales for the test period are about 5 million dollars. The current peanut-butter target says 6.5 million — that is wrong by 30 percent. Our model says 5.14 million — wrong by only 7 percent. So for just this one store, we cut the error by 23 percentage points. And see the green box below? It tells the planner the model is the better choice here. That is one store. Now think about doing this for 1,727 of them."

TIP *If Wolfridge doesn't load, just pick any store that has a green 'Model wins' verdict — the story is the same.*

NEXT Stay on the same store. Click the blue 'Explain' button.

2 • Click EXPLAIN — The GenAI Explainer

1:15 – 1:50 (~35 seconds)

WHAT TO DO

- Click the 'Explain' button next to the heading 'Why did this store get this target?'
- Wait 2 or 3 seconds. A short three-paragraph explanation appears.
- Read the first two lines out loud so the professor can follow.

WHAT TO SAY

"This is our GenAI layer. When the planner clicks Explain, the system writes a plain-English answer for that specific store. It quotes the real numbers — recent sales, year-over-year change, how many competitors are nearby, what the housing market looks like — and tells the planner exactly why the model picked this target and what action to take. Trust it, raise it, or review it. Every one of the 1,727 stores can be explained this way, on demand."

How does this actually work behind the scenes?

When you click Explain, three things happen: **(1)** the app collects that store's 16 feature values along with its predicted target and actual sales. **(2)** It sends this to Google's Gemini AI with a structured prompt asking for a 3-part explanation. **(3)** Gemini writes the paragraph using those exact numbers — never a generic answer.

FALLBACK If Gemini is down or our API quota runs out, the app automatically uses a built-in rule-based explainer that uses the same 16 features and writes a similar 3-paragraph answer. The audience cannot tell the difference. The demo never fails.

TIP If asked 'is this real AI?': YES. Each of the 1,727 stores produces a different explanation because each has different feature values. It is not pre-written text.

NEXT Open the dropdown again and pick an under-targeted store (or move on to the next section as planned).

Bonus • How the GenAI Layer Works

This page is for Q&A — keep it ready in case the professor asks how it works.

The 3-step pipeline (what runs when 'Explain' is clicked)

Step 1: Gather the store's data

The app pulls the 16 feature values for that store from a CSV file — things like sales last week (lag_1), median income, housing share, total competitors, and so on. It also adds the model's predicted target, the peanut-butter target, the actual sales, the bias percentage, and the segment role (High Growth, Defend, etc.).

Step 2: Build the prompt and send to Gemini

The app formats this into a structured prompt asking Gemini for a 3-part answer: (a) Why this target? (b) Why this segment? (c) What action should the planner take? The prompt explicitly tells Gemini to reference the actual numbers — not generic phrases.

Step 3: Gemini writes a custom paragraph

Gemini reads the numbers and writes a unique explanation for that store. Every store gets a different paragraph because every store has different feature values. There is no pre-written template.

The Fallback — why the demo never fails

We built a backup rule-based explainer in Python. It uses the same 16 feature values and writes a similar 3-paragraph answer using if-else logic. If Gemini is unreachable (no internet, API down, quota exhausted), the app silently switches to this fallback. The audience cannot tell the difference. This is important because real production systems need to be robust — they cannot rely on a single external service being available 100 percent of the time.

Is it useful TODAY (with static demo data)?

- Yes. Each of the 1,727 stores has a unique combination of 16 feature values, so the explanation is unique for every store. The data file is fixed, but the model's reasoning is real and store-specific.
- It proves the explainability capability. Planners can interrogate ANY store and understand why the model picked that target — no black box.
- It also serves as a training tool. A new planner can scroll through any store and see how recent sales, income, housing, and competition combine into a target.

How it becomes MORE useful with dynamic, live data

- In production, the data refreshes every week. New actuals come in. Lag features update.

- The same GenAI code runs without any change — Gemini just gets fresh numbers and writes fresh explanations.
- Each week, instead of explaining 1,727 stores, planners only need explanations for the ~91 stores flagged for review. They get an auto-generated briefing per store.
- Over time, the explanations help build organisational memory — planners learn why certain markets behave differently, and the model is no longer 'magic' to them.